

```
In [1]: %load_ext autoreload
        %autoreload 2
        %pylab inline
        %import numpy
        np=numpy
```

Populating the interactive namespace from numpy and matplotlib

```
In [2]: import numpy
import theano
import theano.tensor as T

class MLP(object):
    """multilayer perceptron"""
    def __init__(self, rng, input, n_input, n_hidden, n_output):

        # hidden layer weights
        self.W1 = theano.shared(
            numpy.asarray(
                rng.uniform(
                    low=-numpy.sqrt(1. / n_input),
                    high=numpy.sqrt(1. / n_input),
                    size=(n_input, n_hidden)),
                dtype=theano.config.floatX),
            name='W1', borrow=True)

        # hidden layer biases
        self.b1 = theano.shared(numpy.asarray(numpy.zeros(n_hidden,),
                                                dtype=theano.config.floatX),
                                name='b1', borrow=True)

        # output layer weights
        self.W2 = theano.shared(
            numpy.asarray(
                rng.uniform(
                    low=-numpy.sqrt(1. / n_input),
                    high=numpy.sqrt(1. / n_input),
                    size=(n_hidden, n_output)),
                dtype=theano.config.floatX),
            name='W2', borrow=True)

        # output layer biases
        self.b2 = theano.shared(numpy.asarray(numpy.zeros(n_output,),
                                                dtype=theano.config.floatX),
                                name='b2', borrow=True)

        # prediction formula
        self.p_y_given_x = T.nnet.softmax(T.dot(T.tanh(T.dot(input, self.W1) + self.b1), self.W2) + self.b2)

        # predicted y from x
        self.y_pred = T.argmax(self.p_y_given_x, axis=1)
```

```

        # group all the parameters of the model
        self.params = [self.W1, self.b1, self.W2, self.b2]

        # the penalty functions
        self.L1 = abs(self.W1).sum() + abs(self.W2).sum()
        self.L2 = (self.W1**2).sum() + (self.W2**2).sum()

    def negative_log_likelihood(self, y):
        return -T.mean(T.log(self.p_y_given_x)[T.arange(y.shape[0]), y
1)

    def errors(self, y):
        # check if y has same dimension of y_pred
        if y.ndim != self.y_pred.ndim:
            raise TypeError(
                'y should have the same shape as self.y_pred',
                ('y', y.type, 'y_pred', self.y_pred.type)
            )
        # check if y is of the correct datatype
        if y.dtype.startswith('int'):
            # the T.neq operator returns a vector of 0s and 1s, where
1
            # represents a mistake in prediction
            return T.mean(T.neq(self.y_pred, y))
        else:
            raise NotImplementedError()

```

```

In [3]: def load_data():
        """
        Loads the data
        """

        test_x = numpy.loadtxt('test_images.txt', delimiter=',')
        test_y = numpy.argmax(numpy.loadtxt('test_labels.txt', delimiter=
        ,'), axis=1)
        train_x = numpy.loadtxt('train_images.txt', delimiter=',')
        train_y = numpy.argmax(numpy.loadtxt('train_labels.txt', delimiter
        =','), axis=1)

        def shared_dataset(data_x, data_y, borrow=True):
            """
            Function that loads the dataset into shared variables

            The reason we store our dataset in shared variables is to allo
w
            Theano to copy it into the GPU memory (when code is run on GPU
            ).
            Since copying data into the GPU is slow, copying a minibatch e
verytime
            is needed (the default behaviour if the data is not in a share
d
            variable) would lead to a large decrease in performance.
            """
            shared_x = theano.shared(numpy.asarray(data_x,
                                                    dtype=theano.config.flo
atX),
                                                    borrow=borrow)
            shared_y = theano.shared(numpy.asarray(data_y,
                                                    dtype=theano.config.flo
atX),
                                                    borrow=borrow)
            # When storing data on the GPU it has to be stored as floats
            # therefore we will store the labels as ``floatX`` as well
            # (``shared_y`` does exactly that). But during our computation
s
            # we need them as ints (we use labels as index, and if they ar
e
            # floats it doesn't make sense) therefore instead of returning
            # ``shared_y`` we will have to cast it to int. This little hac
k

            # lets us get around this issue
            return shared_x, T.cast(shared_y, 'int32')

        test_x, test_y = shared_dataset(test_x, test_y)
        train_x, train_y = shared_dataset(train_x, train_y)

        rval = (train_x, train_y, test_x, test_y)
        return rval

```

```

In [4]: def run_test(n_epochs,

```

```

        n_hidden,
        learning_rate=0.01,
        L1_reg=0.00,
        L2_reg=0.0001,
        batch_size=100,
        check_gradients=False):
n_input = 28*28
n_output = 10

# only check the first batch if debug
if check_gradients==True:
    batch_size=1

# load the data
datasets = load_data()
train_set_x, train_set_y, test_set_x, test_set_y = datasets

# compute number of minibatches for training and testing
n_train_batches = train_set_x.get_value(borrow=True).shape[0] / batch_size
n_test_batches = test_set_x.get_value(borrow=True).shape[0] / batch_size

print '\n... building the model'

# allocate symbolic variables for the data
index = T.lscalar() # index to a minibatch
x = T.matrix('x') # the data is presented as rasterized images
y = T.ivector('y') # the labels are presented as 1D vector of integers

rng = numpy.random.RandomState(1234)

# construct the MLP class
classifier = MLP(rng, x, n_input, n_hidden, n_output)

# minimize negative log likelihood of the model
# and the regularization terms (L1 and L2)
# cost is expressed symbolically
cost = (
    classifier.negative_log_likelihood(y)
    + L1_reg * classifier.L1
    + L2_reg * classifier.L2
)

# returns the cost
ret_cost = theano.function(
    inputs=[index],
    outputs=cost,
    givens={
        x: train_set_x[index * batch_size:(index + 1) * batch_size],
        y: train_set_y[index * batch_size:(index + 1) * batch_size]
    }
)

```

```

# fit on train set
check_fit_train_set = theano.function(
    inputs=[],
    outputs=classifier.errors(y),
    givens={
        x: train_set_x,
        y: train_set_y
    }
)

# fit on test set
check_fit_test_set = theano.function(
    inputs=[],
    outputs=classifier.errors(y),
    givens={
        x: test_set_x,
        y: test_set_y
    }
)

# compile a Theano function that computes the mistakes that are made
# by the model on a minibatch of the train set (we'll see overfitting)
check_fit_batch = theano.function(
    inputs=[index],
    outputs=classifier.errors(y),
    givens={
        x: train_set_x[index * batch_size:(index + 1) * batch_size],
        y: train_set_y[index * batch_size:(index + 1) * batch_size]
    }
)

# compute the gradient of cost with respect to theta (sorted in params)
# the resulting gradients will be stored in a list gparams
gradient_w1 = T.grad(cost, classifier.W1)
gradient_b1 = T.grad(cost, classifier.b1)

gradient_w2 = T.grad(cost, classifier.W2)
gradient_b2 = T.grad(cost, classifier.b2)

# specify how to update the parameters of the model as a list of
# (variable, update expression) pairs
updates = [
    (param, param - learning_rate * gradient)
    for param, gradient in [(classifier.W1, gradient_w1),
                           (classifier.b1, gradient_b1),
                           (classifier.W2, gradient_w2),
                           (classifier.b2, gradient_b2)]
]

values = [gradient_w1, gradient_b1, gradient_w2, gradient_b2]

ret_gradient = theano.function(

```

```

        inputs=[index],
        outputs=values,
        givens={
            x: train_set_x[index * batch_size: (index + 1) * batch_size],
            y: train_set_y[index * batch_size: (index + 1) * batch_size]
        }
    )

    # updates the parameter of the model based on the rules defined in
    'updates'
    train_model = theano.function(
        inputs=[index],
        outputs=[],
        updates=updates,
        givens={
            x: train_set_x[index * batch_size: (index + 1) * batch_size],
            y: train_set_y[index * batch_size: (index + 1) * batch_size]
        }
    )

    print '... training'

    best_iter = 0
    test_scores = []
    train_scores = []

    epoch = 0

    while (epoch < n_epochs):

        epoch = epoch + 1

        for minibatch_index in xrange(n_train_batches):
            train_model(minibatch_index)
            if check_gradients==True:
                symbolic_gradients = ret_gradient(minibatch_index)

                epsilon = 1E-4
                a = numpy.zeros((n_input, n_hidden))
                b = numpy.zeros((n_hidden,))
                c = numpy.zeros((n_hidden, n_output))
                d = numpy.zeros((n_output,))
                grad= [a,b,c,d]

                # b1
                b1_vals = classifier.b1.get_value()
                for i in range(len(grad[1])):
                    # compute low
                    b1_vals[i] -= epsilon
                    classifier.b1.set_value(b1_vals, borrow=True)
                    low= ret_cost(minibatch_index)

```

```

# compute high
b1_vals[i] += 2.*epsilon
classifier.b1.set_value(b1_vals, borrow=True)
high= ret_cost(minibatch_index)

# reset the value
b1_vals[i] -= epsilon
# store the gradient
grad[1][i] = (high - low) / (2.*epsilon)

# b2
b2_vals = classifier.b2.get_value()
for i in range(len(grad[3])):
    # compute low
    b2_vals[i] -= epsilon
    classifier.b2.set_value(b2_vals, borrow=True)
    low= ret_cost(minibatch_index)

    # compute high
    b2_vals[i] += 2.*epsilon
    classifier.b2.set_value(b2_vals, borrow=True)
    high= ret_cost(minibatch_index)

    # reset the value
    b2_vals[i] -= epsilon

    # store the gradient
    grad[3][i] = (high - low) / (2.*epsilon)

# W1
w1_vals = classifier.W1.get_value()
for i in range(len(grad[0])):
    for j in range(len(grad[0][0])):

        # compute low
        w1_vals[i][j] -= epsilon
        classifier.W1.set_value(w1_vals, borrow=True)
        low= ret_cost(minibatch_index)

        # compute high
        w1_vals[i][j] += 2.*epsilon
        classifier.W1.set_value(w1_vals, borrow=True)
        high= ret_cost(minibatch_index)

        # reset the value
        w1_vals[i][j] -= epsilon

        # store the gradient
        grad[0][i][j] = (high - low) / (2.*epsilon)

# W2
w2_vals = classifier.W2.get_value()
for i in range(len(grad[2])):
    for j in range(len(grad[2][0])):

```

```

        # compute low
        w2_vals[i][j] -= epsilon
        classifier.W2.set_value(w2_vals, borrow=True)
        low= ret_cost(minibatch_index)

        # compute high
        w2_vals[i][j] += 2.*epsilon
        classifier.W2.set_value(w2_vals, borrow=True)
        high= ret_cost(minibatch_index)

        # reset the value
        w2_vals[i][j] -= epsilon

        # store the gradient
        grad[2][i][j] = (high - low) / (2.*epsilon)

    return (symbolic_gradients, grad)

# verify the fit on the datasets
train_scores.append(check_fit_train_set())
test_scores.append(check_fit_test_set())

print '... done\n'
return (train_scores, test_scores)

```

```
In [5]: train_scores, test_scores = run_test(1000, 500)
```

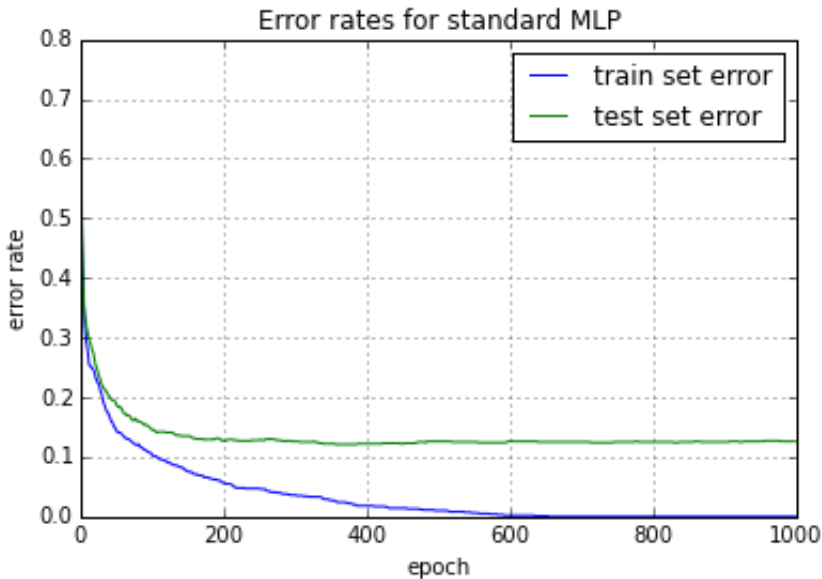
```

... building the model
... training
... done

```



```
In [6]: plt.plot(train_scores, label='train set error')
plt.plot(test_scores, label='test set error')
plt.xlabel('epoch')
plt.ylabel('error rate')
plt.grid(True)
plt.legend()
plt.title('Error rates for standard MLP')
plt.show()
print "Parameters are the following: epochs=1000, n_hidden=500, learning_rate=0.01, L1_reg=0.00, L2_reg=0.0001, batch_size=100"
```



Parameters are the following: epochs=1000, n_hidden=500, learning_rate=0.01, L1_reg=0.00, L2_reg=0.0001, batch_size=100

```
In [8]: import matplotlib.pyplot as plt

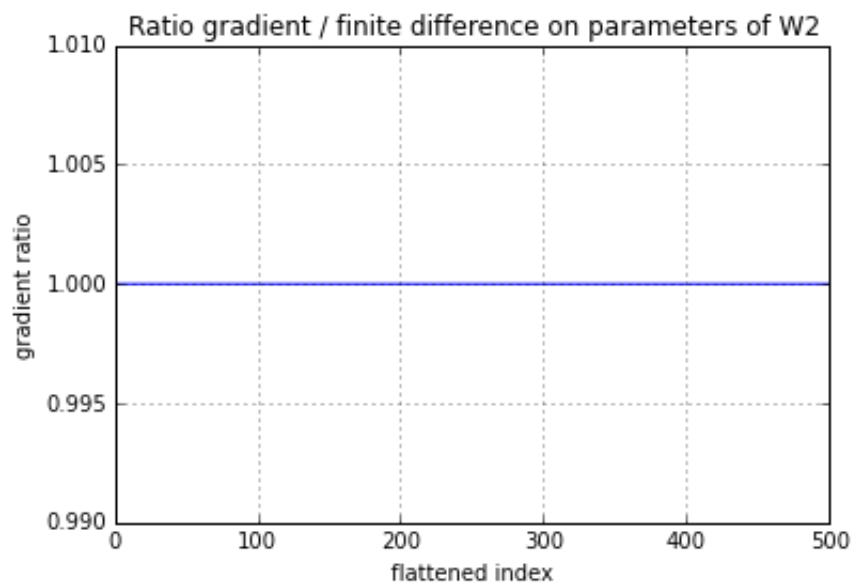
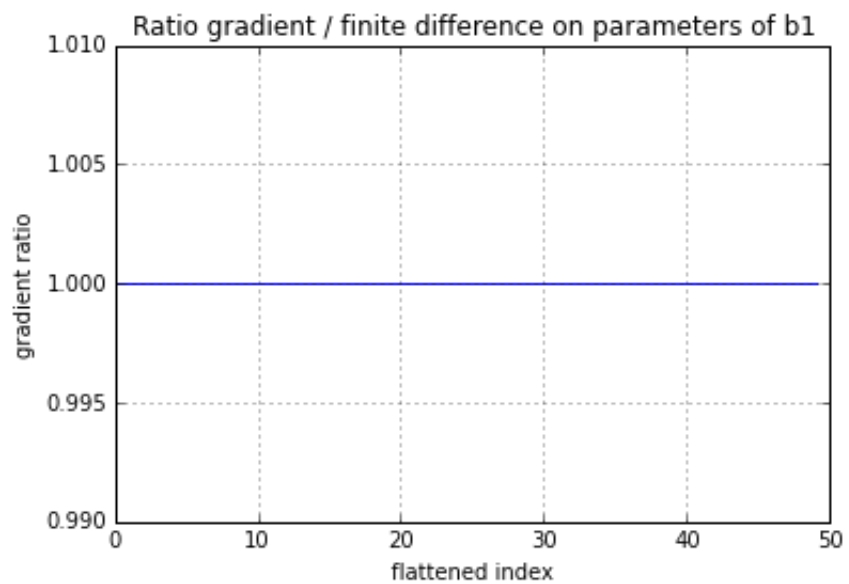
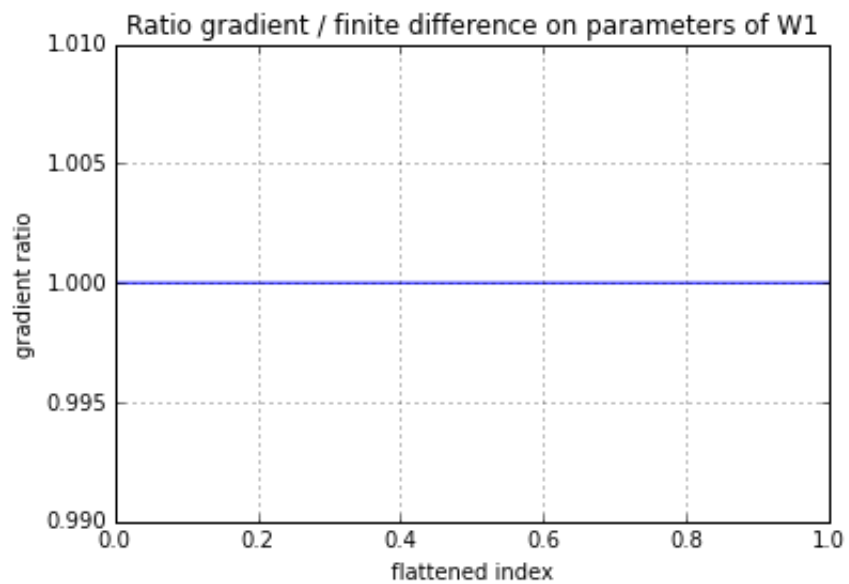
theano_grad, finite_diff_grad = run_test(1000, 50, check_gradients=True)

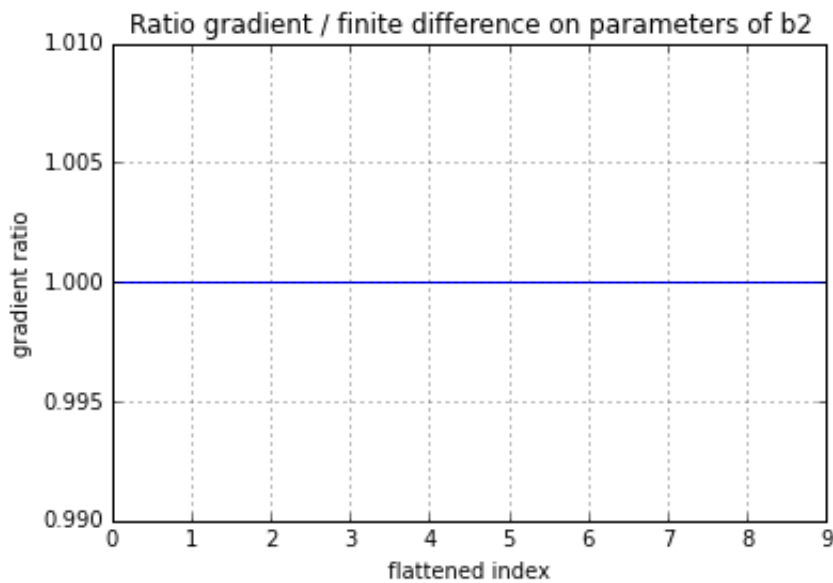
pylab.gca().set_autoscale_on(False)

theano_grad_flat = map(numpy.ndarray.flatten, theano_grad)
finite_diff_grad_flat = map(numpy.ndarray.flatten, finite_diff_grad)

labels = ['W1', 'b1', 'W2', 'b2']
for i in range(len(theano_grad_flat)):
    plt.figure(i+1)
    data = theano_grad_flat[i] / finite_diff_grad_flat[i]
    plt.plot(data)
    plt.xlabel('flattened index')
    plt.ylabel('gradient ratio')
    plt.grid(True)
    plt.title('Ratio gradient / finite difference on parameters of {0}'
              .format(labels[i]))
    ylim((0.99, 1.01))
```

```
... building the model
... training
```





```
In [9]: class MLP_relu(MLP):
        # modified with max(0, ha)
        def __init__(self, rng, input, n_input, n_hidden, n_output):
            super(MLP_relu, self).__init__(rng, input, n_input, n_hidden,
            n_output)
            self.p_y_given_x = T.nnet.softmax(T.dot(T.tanh(T.maximum(T.dot
            (input, self.W1) + self.b1, 0.)), self.W2) + self.b2)
```

```
In [10]: def run_test_relu(n_epochs,
                        n_hidden,
                        learning_rate=0.01,
                        L1_reg=0.00,
                        L2_reg=0.0001,
                        batch_size=100,
                        check_gradients=False):

    n_input = 28*28
    n_output = 10

    # only check the first batch if debug
    if check_gradients==True:
        batch_size=1

    # load the data
    datasets = load_data()
    train_set_x, train_set_y, test_set_x, test_set_y = datasets

    # compute number of minibatches for training and testing
    n_train_batches = train_set_x.get_value(borrow=True).shape[0] / batch_size
    n_test_batches = test_set_x.get_value(borrow=True).shape[0] / batch_size

    print '\n... building the model'

    # allocate symbolic variables for the data
    index = T.lscalar() # index to a minibatch
    x = T.matrix('x') # the data is presented as rasterized images
```

```

    y = T.ivector('y')    # the labels are presented as 1D vector of in
t labels

    rng = numpy.random.RandomState(1234)

    # construct the MLP class
    classifier = MLP_relu(rng, x, n_input, n_hidden, n_output)

    # minimize negative log likelihood of the model
    # and the regularization terms (L1 and L2)
    # cost is expressed symbolically
    cost = (
        classifier.negative_log_likelihood(y)
        + L1_reg * classifier.L1
        + L2_reg * classifier.L2
    )

    # returns the cost
    ret_cost = theano.function(
        inputs=[index],
        outputs=cost,
        givens={
            x: train_set_x[index * batch_size:(index + 1) * batch_size
],
            y: train_set_y[index * batch_size:(index + 1) * batch_size
]
        }
    )

    # fit on train set
    check_fit_train_set = theano.function(
        inputs=[],
        outputs=classifier.errors(y),
        givens={
            x: train_set_x,
            y: train_set_y
        }
    )

    # fit on test set
    check_fit_test_set = theano.function(
        inputs=[],
        outputs=classifier.errors(y),
        givens={
            x: test_set_x,
            y: test_set_y
        }
    )

    # compile a Theano function that computes the mistakes that are ma
de
    # by the model on a minibatch of the train set (we'll see overfitt
ing)
    check_fit_batch = theano.function(
        inputs=[index],
        outputs=classifier.errors(y),
        givens={

```

```

        x: train_set_x[index * batch_size:(index + 1) * batch_size
    ],
        y: train_set_y[index * batch_size:(index + 1) * batch_size
    ]
    }
    )

    # compute the gradient of cost with respect to theta (sorted in pa
rams)
    # the resulting gradients will be stored in a list gparams
    gradient_w1 = T.grad(cost, classifier.W1)
    gradient_b1 = T.grad(cost, classifier.b1)

    gradient_w2 = T.grad(cost, classifier.W2)
    gradient_b2 = T.grad(cost, classifier.b2)

    # specify how to update the parameters of the model as a list of
    # (variable, update expression) pairs
    updates = [
        (param, param - learning_rate * gradient)
        for param, gradient in [(classifier.W1, gradient_w1),
                                (classifier.b1, gradient_b1),
                                (classifier.W2, gradient_w2),
                                (classifier.b2, gradient_b2)]
    ]

    values = [gradient_w1, gradient_b1, gradient_w2, gradient_b2]

    ret_gradient = theano.function(
        inputs=[index],
        outputs=values,
        givens={
            x: train_set_x[index * batch_size: (index + 1) * batch_siz
e],
            y: train_set_y[index * batch_size: (index + 1) * batch_siz
e]
        }
    )

    # updates the parameter of the model based on the rules defined in
    'updates'
    train_model = theano.function(
        inputs=[index],
        outputs=[],
        updates=updates,
        givens={
            x: train_set_x[index * batch_size: (index + 1) * batch_siz
e],
            y: train_set_y[index * batch_size: (index + 1) * batch_siz
e]
        }
    )

    print '... training'

    best_iter = 0
    test_scores = []

```

```

train_scores = []

epoch = 0

while (epoch < n_epochs):

    epoch = epoch + 1

    for minibatch_index in xrange(n_train_batches):
        train_model(minibatch_index)
        if check_gradients==True:
            symbolic_gradients = ret_gradient(minibatch_index)

            epsilon = 1E-4
            a = numpy.zeros((n_input, n_hidden))
            b = numpy.zeros((n_hidden,))
            c = numpy.zeros((n_hidden, n_output))
            d = numpy.zeros((n_output,))
            grad= [a,b,c,d]

            # b1
            b1_vals = classifier.b1.get_value()
            for i in range(len(grad[1])):
                # compute low
                b1_vals[i] -= epsilon
                classifier.b1.set_value(b1_vals, borrow=True)
                low= ret_cost(minibatch_index)

                # compute high
                b1_vals[i] += 2.*epsilon
                classifier.b1.set_value(b1_vals, borrow=True)
                high= ret_cost(minibatch_index)

                # reset the value
                b1_vals[i] -= epsilon
                # store the gradient
                grad[1][i] = (high - low) / (2.*epsilon)

            # b2
            b2_vals = classifier.b2.get_value()
            for i in range(len(grad[3])):
                # compute low
                b2_vals[i] -= epsilon
                classifier.b2.set_value(b2_vals, borrow=True)
                low= ret_cost(minibatch_index)

                # compute high
                b2_vals[i] += 2.*epsilon
                classifier.b2.set_value(b2_vals, borrow=True)
                high= ret_cost(minibatch_index)

                # reset the value
                b2_vals[i] -= epsilon

                # store the gradient

```

```

        grad[3][i] = (high - low) / (2.*epsilon)

# W1
w1_vals = classifier.W1.get_value()
for i in range(len(grad[0])):
    for j in range(len(grad[0][0])):

        # compute low
        w1_vals[i][j] -= epsilon
        classifier.W1.set_value(w1_vals, borrow=True)
        low = ret_cost(minibatch_index)

        # compute high
        w1_vals[i][j] += 2.*epsilon
        classifier.W1.set_value(w1_vals, borrow=True)
        high = ret_cost(minibatch_index)

        # reset the value
        w1_vals[i][j] -= epsilon

        # store the gradient
        grad[0][i][j] = (high - low) / (2.*epsilon)

# W2
w2_vals = classifier.W2.get_value()
for i in range(len(grad[2])):
    for j in range(len(grad[2][0])):

        # compute low
        w2_vals[i][j] -= epsilon
        classifier.W2.set_value(w2_vals, borrow=True)
        low = ret_cost(minibatch_index)

        # compute high
        w2_vals[i][j] += 2.*epsilon
        classifier.W2.set_value(w2_vals, borrow=True)
        high = ret_cost(minibatch_index)

        # reset the value
        w2_vals[i][j] -= epsilon

        # store the gradient
        grad[2][i][j] = (high - low) / (2.*epsilon)

    return (symbolic_gradients, grad)

# verify the fit on the datasets
train_scores.append(check_fit_train_set())
test_scores.append(check_fit_test_set())

print '... done\n'
return (train_scores, test_scores)

```

```
In [11]: train_scores, test_scores = run_test_relu(1000, 500)
```

```
... building the model  
... training  
... done
```

```
In [12]: plt.plot(train_scores, label='train set error')  
plt.plot(test_scores, label='test set error')  
plt.xlabel('epoch')  
plt.ylabel('error rate')  
plt.grid(True)  
plt.legend()  
plt.title('Error rates for RELU (500 hidden,1000 epochs)')  
plt.show()
```

